

```
1 import os
2 import pandas as pd
3 import numpy as np
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6 import tensorflow as tf
7 from tensorflow.keras.models import Model, load_model
8 from tensorflow.keras.layers import (Input, Conv2D, MaxPooling2D,
9                                         Flatten, Dense, Dropout)
10 from tensorflow.keras.optimizers import Adam
11 from tensorflow.keras.preprocessing.image import ImageDataGenerator
12 from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau
13 from sklearn.model_selection import train_test_split
14 from sklearn.metrics import classification_report, confusion_matrix, Conf
```

```
1 !mkdir -p ~/.kaggle
2 !cp kaggle.json ~/.kaggle/
3 !chmod 600 ~/.kaggle/kaggle.json
4
```

```
1 !pip install --upgrade kaggle
2 !kaggle --version
3
```

→ Requirement already satisfied: kaggle in /usr/local/lib/python3.11/dist-packages (1.7.4).
Requirement already satisfied: bleach in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: certifi>=14.05.14 in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: charset-normalizer in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: idna in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: protobuf in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: python-dateutil>=2.5.3 in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: python-slugify in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: setuptools>=21.0.0 in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: six>=1.10 in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: text-unidecode in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: urllib3>=1.15.1 in /usr/local/lib/python3.11/dist-packages (from kaggle)
Requirement already satisfied: webencodings in /usr/local/lib/python3.11/dist-packages (from kaggle)
Kaggle API 1.7.4.2

```
1 !kaggle datasets download -d masoudnickparvar/brain-tumor-mri-dataset -p /content
```

```
2
```

→ Dataset URL: <https://www.kaggle.com/datasets/masoudnickparvar/brain-tumor-mri-dataset>
License(s): CC0-1.0

```

1 def get_class_paths(base_path):
2     classes, paths = [], []
3     for label in os.listdir(base_path):
4         full_dir = os.path.join(base_path, label)
5         if os.path.isdir(full_dir):
6             for fname in os.listdir(full_dir):
7                 classes.append(label)
8                 paths.append(os.path.join(full_dir, fname))
9     return pd.DataFrame({"Class": classes, "Path": paths})
10
11 tr_df = get_class_paths("Training")
12 ts_df = get_class_paths("Testing")

1 #Split Testing into validation & test 50/50
2 valid_df, ts_df = train_test_split(
3     ts_df, stratify=ts_df["Class"], train_size=0.5, random_state=42
4 )
5 #RANDOM STATE Sets a seed for reproducibility
6 #ensuring that the split is consistent across different runs
7 #every time run the code
8 # the data will be split in the same way ensuring consistency in result
9 batch_size = 16
10 img_size = (224, 224)
11
12 train_datagen = ImageDataGenerator(rescale=1/255, brightness_range=(0.8,1.2))
13 test_datagen = ImageDataGenerator(rescale=1/255)
14
15 tr_gen = train_datagen.flow_from_dataframe(
16     tr_df, x_col="Path", y_col="Class",
17     target_size=img_size, batch_size=batch_size
18 )
19 valid_gen = train_datagen.flow_from_dataframe(
20     valid_df, x_col="Path", y_col="Class",
21     target_size=img_size, batch_size=batch_size
22 )
23 ts_gen = test_datagen.flow_from_dataframe(
24     ts_df, x_col="Path", y_col="Class",
25     target_size=img_size, batch_size=batch_size, shuffle=False
26 )

```

→ Found 5712 validated image filenames belonging to 4 classes.
 Found 655 validated image filenames belonging to 4 classes.
 Found 656 validated image filenames belonging to 4 classes.

```

1 #VGG-19 Architecture
2
3 inputs = Input(shape=(224, 224, 3))
4

```

```
5 # Block 1 (2 Conv64 + Pool)
6 x = Conv2D(64, (3,3), activation="relu", padding="same")(inputs)
7 x = Conv2D(64, (3,3), activation="relu", padding="same")(x)
8 x = MaxPooling2D((2,2), strides=2)(x)
9
10 # Block 2 (2 Conv128 + Pool)
11 x = Conv2D(128, (3,3), activation="relu", padding="same")(x)
12 x = Conv2D(128, (3,3), activation="relu", padding="same")(x)
13 x = MaxPooling2D((2,2), strides=2)(x)
14
15 # Block 3 (4 Conv256 + Pool)
16 for _ in range(4):
17     x = Conv2D(256, (3,3), activation="relu", padding="same")(x)
18 x = MaxPooling2D((2,2), strides=2)(x)
19
20 # Block 4 (4 Conv512 + Pool)
21 for _ in range(4):
22     x = Conv2D(512, (3,3), activation="relu", padding="same")(x)
23 x = MaxPooling2D((2,2), strides=2)(x)
24
25 # Block 5 (4 Conv512 + Pool)
26 for _ in range(4):
27     x = Conv2D(512, (3,3), activation="relu", padding="same")(x)
28 x = MaxPooling2D((2,2), strides=2)(x)
29
30
```

```
1 # flatten layer
2 x = Flatten()(x)
3 x = Dense(4096, activation="relu")(x)
4 x = Dense(4096, activation="relu")(x) #Stacking dense layers enables the model to
5
6
7 outputs = Dense(len(tr_gen.class_indices), activation="softmax")(x)
8
9 model = Model(inputs, outputs, name="VGG19_Custom")
10 model.compile(
11     optimizer=Adam(learning_rate=1e-4),#computes adaptive learning rates for each
12     loss="categorical_crossentropy",
13     metrics=["accuracy"]
14 )
15 model.summary()
16
```


 Model: "VGG19_Custom"

Layer (type)	Output Shape	Param #
input_layer (InputLayer)	(None, 224, 224, 3)	0
conv2d (Conv2D)	(None, 224, 224, 64)	1,792
conv2d_1 (Conv2D)	(None, 224, 224, 64)	36,928
max_pooling2d (MaxPooling2D)	(None, 112, 112, 64)	0
conv2d_2 (Conv2D)	(None, 112, 112, 128)	73,856
conv2d_3 (Conv2D)	(None, 112, 112, 128)	147,584
max_pooling2d_1 (MaxPooling2D)	(None, 56, 56, 128)	0
conv2d_4 (Conv2D)	(None, 56, 56, 256)	295,168
conv2d_5 (Conv2D)	(None, 56, 56, 256)	590,080
conv2d_6 (Conv2D)	(None, 56, 56, 256)	590,080
conv2d_7 (Conv2D)	(None, 56, 56, 256)	590,080
max_pooling2d_2 (MaxPooling2D)	(None, 28, 28, 256)	0
conv2d_8 (Conv2D)	(None, 28, 28, 512)	1,180,160
conv2d_9 (Conv2D)	(None, 28, 28, 512)	2,359,808
conv2d_10 (Conv2D)	(None, 28, 28, 512)	2,359,808
conv2d_11 (Conv2D)	(None, 28, 28, 512)	2,359,808
max_pooling2d_3 (MaxPooling2D)	(None, 14, 14, 512)	0
conv2d_12 (Conv2D)	(None, 14, 14, 512)	2,359,808
conv2d_13 (Conv2D)	(None, 14, 14, 512)	2,359,808
conv2d_14 (Conv2D)	(None, 14, 14, 512)	2,359,808
conv2d_15 (Conv2D)	(None, 14, 14, 512)	2,359,808
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 4096)	102,764,544
dense_1 (Dense)	(None, 4096)	16,781,312
dense_2 (Dense)	(None, 4)	16,388

```
1 #Train model
2
3 early_stop  = EarlyStopping(monitor="val_loss", patience=3, restore_best_weights=True)
4 #EarlyStopping Callback set to monitor the val_loss if the validation loss doesn't
5 reduce_lr    = ReduceLROnPlateau(monitor="val_loss", factor=0.5, patience=2, verbose=1)
6
7 history = model.fit(
8     tr_gen,
9     epochs=20,
10    validation_data=valid_gen,
11    callbacks=[early_stop, reduce_lr]
12 )
```


```
→ /usr/local/lib/python3.11/dist-packages/keras/src/trainers/data_adapters/py_dataset_adapter.py:115: UserWarning: Model training is taking longer than expected. This may be due to a slow disk or network. If you are training on a local disk, consider using a solid-state drive or an SSD. If you are training on a network, consider using a faster network or a faster server. You can also try increasing the batch size or reducing the number of epochs.
  self._warn_if_super_not_called()
Epoch 1/20
357/357 [██████████] 182s 400ms/step - accuracy: 0.4367 - loss: 1.1712 - val_accuracy: 0.4367 - val_loss: 1.1712
Epoch 2/20
357/357 [██████████] 129s 361ms/step - accuracy: 0.7414 - loss: 0.6064 - val_accuracy: 0.7414 - val_loss: 0.6064
Epoch 3/20
357/357 [██████████] 129s 361ms/step - accuracy: 0.7981 - loss: 0.4822 - val_accuracy: 0.7981 - val_loss: 0.4822
Epoch 4/20
357/357 [██████████] 129s 361ms/step - accuracy: 0.8532 - loss: 0.3793 - val_accuracy: 0.8532 - val_loss: 0.3793
Epoch 5/20
357/357 [██████████] 142s 361ms/step - accuracy: 0.8897 - loss: 0.2988 - val_accuracy: 0.8897 - val_loss: 0.2988
Epoch 6/20
357/357 [██████████] 129s 360ms/step - accuracy: 0.9116 - loss: 0.2332 - val_accuracy: 0.9116 - val_loss: 0.2332
Epoch 7/20
357/357 [██████████] 129s 360ms/step - accuracy: 0.9304 - loss: 0.1821 - val_accuracy: 0.9304 - val_loss: 0.1821
Epoch 8/20
357/357 [██████████] 129s 361ms/step - accuracy: 0.9448 - loss: 0.1559 - val_accuracy: 0.9448 - val_loss: 0.1559
Epoch 9/20
357/357 [██████████] 128s 359ms/step - accuracy: 0.9567 - loss: 0.1206 - val_accuracy: 0.9567 - val_loss: 0.1206
Epoch 10/20
357/357 [██████████] 0s 344ms/step - accuracy: 0.9517 - loss: 0.1300
Epoch 10: ReduceLROnPlateau reducing learning rate to 4.999999873689376e-05.
357/357 [██████████] 128s 359ms/step - accuracy: 0.9517 - loss: 0.1300 - val_accuracy: 0.9517 - val_loss: 0.1300
Epoch 11/20
357/357 [██████████] 128s 360ms/step - accuracy: 0.9757 - loss: 0.0692 - val_accuracy: 0.9757 - val_loss: 0.0692
Epoch 12/20
357/357 [██████████] 129s 360ms/step - accuracy: 0.9877 - loss: 0.0347 - val_accuracy: 0.9877 - val_loss: 0.0347
Epoch 13/20
357/357 [██████████] 129s 360ms/step - accuracy: 0.9927 - loss: 0.0250 - val_accuracy: 0.9927 - val_loss: 0.0250
Epoch 14/20
357/357 [██████████] 141s 358ms/step - accuracy: 0.9934 - loss: 0.0179 - val_accuracy: 0.9934 - val_loss: 0.0179
Epoch 15/20
357/357 [██████████] 0s 344ms/step - accuracy: 0.9986 - loss: 0.0069
Epoch 15: ReduceLROnPlateau reducing learning rate to 2.499999936844688e-05.
357/357 [██████████] 128s 358ms/step - accuracy: 0.9986 - loss: 0.0069 - val_accuracy: 0.9986 - val_loss: 0.0069
Epoch 16/20
357/357 [██████████] 128s 359ms/step - accuracy: 0.9979 - loss: 0.0052 - val_accuracy: 0.9979 - val_loss: 0.0052
```

```
1 # Plot accuracy loss
2 plt.figure(figsize=(12,5))
3
4 # Accuracy
5 plt.subplot(1,2,1)
6 plt.plot(history.history["accuracy"], label="Train Acc")
7 plt.plot(history.history["val_accuracy"], label="Val Acc")
8 plt.title("Accuracy over Epochs")
9 plt.xlabel("Epoch")
10 plt.ylabel("Accuracy")
11 plt.legend()
12 plt.grid(True)
13
14
```



```
1 # Loss
2 plt.subplot(1,2,2)
3 plt.plot(history.history["loss"], label="Train Loss")
4 plt.plot(history.history["val_loss"], label="Val Loss")
5 plt.title("Loss over Epochs")
6 plt.xlabel("Epoch")
7 plt.ylabel("Loss")
8 plt.legend()
9 plt.grid(True)
```

```
10  
11 plt.tight_layout()  
12 plt.show()  
13  
14 model.save("vgg19_brain_tumor_mri.h5")  
15
```



WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.savi

```
1 # Evaluate on test set  
2  
3 # predictions  
4 y_true = ts_gen.classes  
5 y_pred = np.argmax(model.predict(ts_gen), axis=1)  
6  
7 # confusion Matrix  
8 cm = confusion_matrix(y_true, y_pred)  
9 disp = ConfusionMatrixDisplay(cm, display_labels=list(tr_gen.class_indices.keys()))  
10  
11 plt.figure(figsize=(7,6))  
12 disp.plot(cmap="Blues", xticks_rotation=45)  
13 plt.title("Test Set Confusion Matrix")  
14 plt.grid(False)  
15 plt.show()
```

16


17 print("Classification Report:\n")

18 print(classification_report(y_true, y_pred, target_names=list(tr_gen.class_indices)))

19

41/41 ————— 6s 124ms/step

<Figure size 700x600 with 0 Axes>



Classification Report:

	precision	recall	f1-score	support
glioma	0.97	0.94	0.96	150
meningioma	0.95	0.95	0.95	153
notumor	0.98	1.00	0.99	203
pituitary	0.99	1.00	1.00	150
accuracy			0.97	656
macro avg	0.97	0.97	0.97	656
weighted avg	0.97	0.97	0.97	656

1 #DISCUSSION

2 The role of each layer in the architecture

3 Input layer: Accepts images of size 224x224x3

4 CONV LAYER: 3x3 filters to the input image, detecting various features such as edges and textures.

5 ReLU activation: After each convolutional operation ReLU activation funct

6 Max pool:Max pooling involves sliding a fixed-size window 2x2 over the ir

7 Fully connected layer: At the end of the network fully connected layers i

8 Softmax layer:Softmax layer outputs a probability distribution over the r

9

10

11

12 Why VGG-19 uses small filters (3x3) [1]

13 Small filters allow the network to capture detailed spatial hierarchies i

1 Start coding or generate with AI.